

Language Models (and representations)

Yoav Goldberg

Last Time

- Multi-Layer perceptrons

$$f_{\theta}(\mathbf{x}) = \text{NN}_{\text{MLP2}}(\mathbf{x}) = \mathbf{y}$$

$$\mathbf{h}^1 = g^1(\mathbf{x}\mathbf{W}^1 + \mathbf{b}^1)$$

$$\mathbf{h}^2 = g^2(\mathbf{h}^1\mathbf{W}^2 + \mathbf{b}^2)$$

$$\mathbf{y} = \mathbf{h}^2\mathbf{W}^3$$

Let's talk about sequences

- Predicting how a sequence will continue.

תל א_

ראיתי כלב חז

מה אתה _

Language Model

$$p(x_i | x_1, \dots, x_{i-1})$$

Language Model: Markov Assumption

$$p(x_i | x_1, \dots, x_{i-1}) \approx p(x_i | x_{i-4}, x_{i-3}, x_{i-2}, x_{i-1})$$

Language Model: Markov Assumption

$$p(x_i | x_1, \dots, x_{i-1}) \approx p(x_i | x_{i-4}, x_{i-3}, x_{i-2}, x_{i-1})$$

(condition only on last n items)

this is called n-gram language model

Language Model

- LM can also be used to assign a probability to a sequence.

$$\begin{aligned} p(x_1, \dots, x_n) = & p_{LM}(x_1 | *S*, *S*) \\ & \times p_{LM}(x_2 | *S*, x_1) \\ & \times p_{LM}(x_3 | x_1, x_2) \\ & \times p_{LM}(x_4 | x_2, x_3) \\ & \dots \\ & \times p_{LM}(x_n | x_{n-2}, x_{n-1}) \end{aligned}$$

Language Model

- Very useful (used in Speech Recognition, Machine Translation.. and many others).
- Does not have to be over natural language.
- Huge research topic. We'll see a neural LM.

Neural LM

$$p(x_k | x_{k-4}, x_{k-3}, x_{k-2}, x_{k-1}) = \text{softmax}(\text{MLP}(\mathbf{x}))$$

$$\mathbf{x} = \text{encode}(x_{k-4}, x_{k-3}, x_{k-2}, x_{k-1})$$

Neural LM

$$p(x_k | x_{k-4}, x_{k-3}, x_{k-2}, x_{k-1}) = \text{softmax}(\text{MLP}(\mathbf{x}))$$

$$\mathbf{x} = \text{encode}(x_{k-4}, x_{k-3}, x_{k-2}, x_{k-1})$$

$$\text{softmax}(g(g(\mathbf{x}\mathbf{W}^1 + \mathbf{b}^1)\mathbf{W}^2 + \mathbf{b}^2)\mathbf{W}^3 + \mathbf{b}^3)$$

$\text{softmax}(\square)$

\uparrow

$\square \mathbf{W}^3 + \mathbf{b}^3$

\uparrow

$g(\square \mathbf{W}^2 + \mathbf{b}^2)$

\uparrow

$g(\square \mathbf{W}^1 + \mathbf{b}^1)$

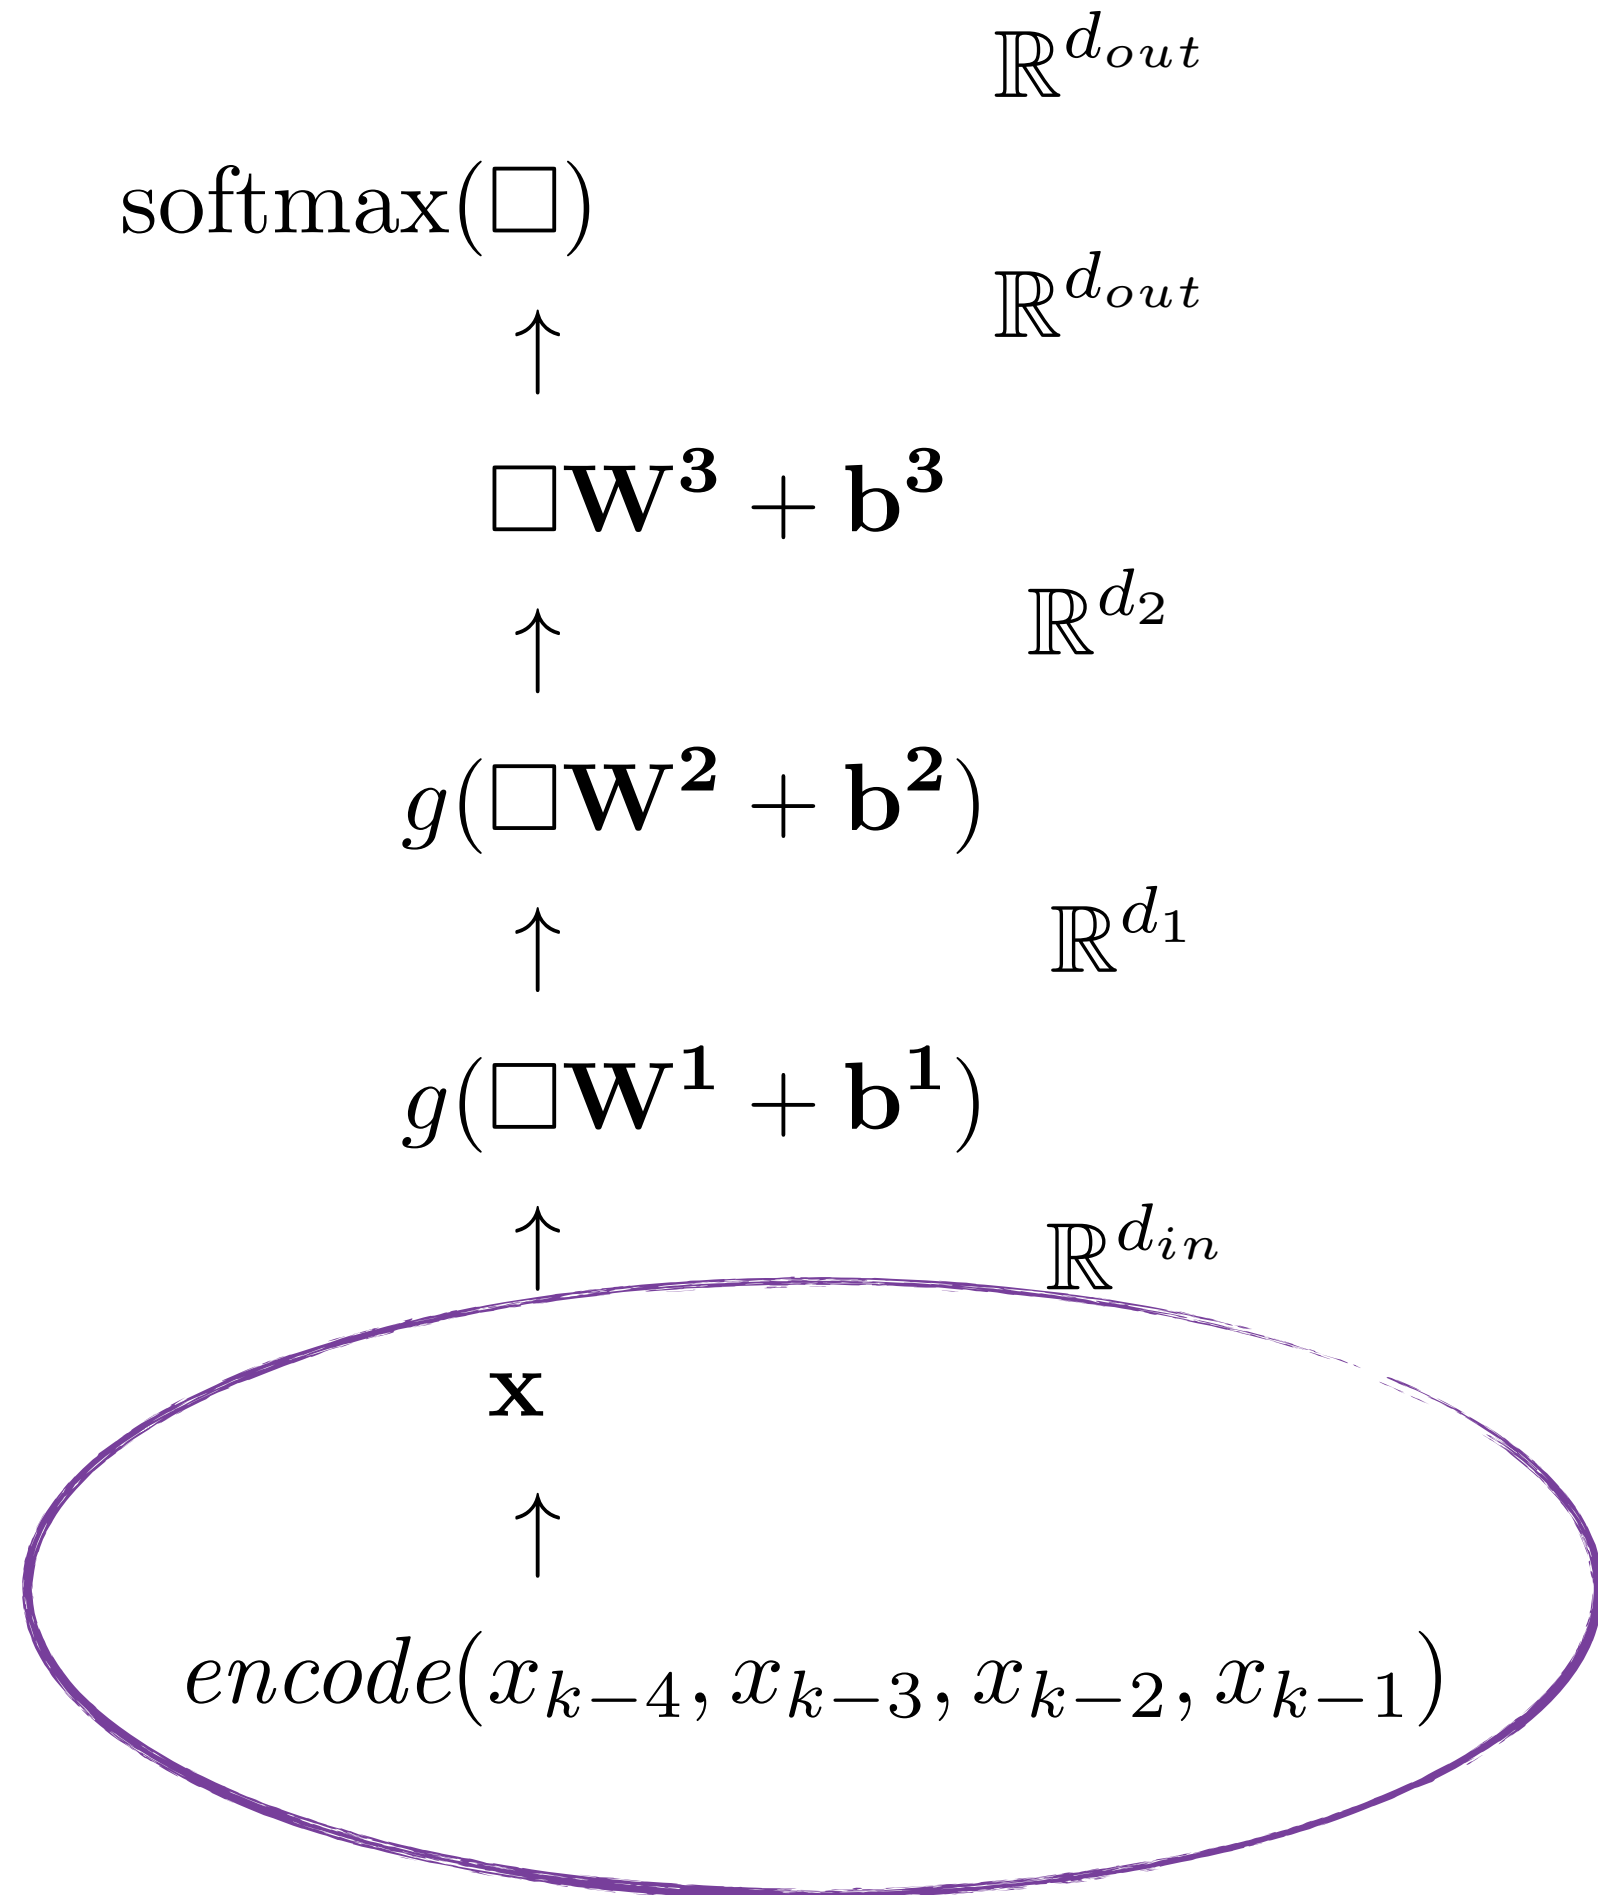
\uparrow

\mathbf{x}

\uparrow

$encode(x_{k-4}, x_{k-3}, x_{k-2}, x_{k-1})$

$$\begin{array}{c}
\mathbb{R}^{d_{out}} \\
\text{softmax}(\square) \\
\uparrow \\
\square \mathbf{W}^3 + \mathbf{b}^3 \\
\uparrow \\
\mathbb{R}^{d_2} \\
g(\square \mathbf{W}^2 + \mathbf{b}^2) \\
\uparrow \\
\mathbb{R}^{d_1} \\
g(\square \mathbf{W}^1 + \mathbf{b}^1) \\
\uparrow \\
\mathbb{R}^{d_{in}} \\
\mathbf{x} \\
\uparrow \\
\text{encode}(x_{k-4}, x_{k-3}, x_{k-2}, x_{k-1})
\end{array}$$



Encoding k elements

$$\text{encode}(x_1, x_2, x_3, x_4)$$

We have k elements in a vocabulary of size $|V|$

Encoding k elements

$$\text{encode}(x_1, x_2, x_3, x_4)$$

We have k elements in a vocabulary of size |V|

4

10

$$V=\{A,B,C,D,E,F,G,H,I,J\}$$

Encoding k elements

$\text{encode}(D, A, G, C)$

A= [1,0,0,0,0,0,0,0,0,0]

B= [0,1,0,0,0,0,0,0,0,0]

C= [0,0,1,0,0,0,0,0,0,0]

D= [0,0,0,1,0,0,0,0,0,0]

E= [0,0,0,0,1,0,0,0,0,0]

F= [0,0,0,0,0,1,0,0,0,0]

G= [0,0,0,0,0,0,1,0,0,0]

H= [0,0,0,0,0,0,0,1,0,0]

I= [0,0,0,0,0,0,0,0,1,0]

J= [0,0,0,0,0,0,0,0,0,1]

Encoding k elements

$\text{encode}(D, A, G, C)$

A= [1,0,0,0,0,0,0,0,0,0]

B= [0,1,0,0,0,0,0,0,0,0]

C= [0,0,1,0,0,0,0,0,0,0]

D= [0,0,0,1,0,0,0,0,0,0]

E= [0,0,0,0,1,0,0,0,0,0]

F= [0,0,0,0,0,1,0,0,0,0]

G= [0,0,0,0,0,0,1,0,0,0]

H= [0,0,0,0,0,0,0,1,0,0]

I= [0,0,0,0,0,0,0,0,1,0]

J= [0,0,0,0,0,0,0,0,0,1]

$\mathbf{v}_D + \mathbf{v}_A + \mathbf{v}_G + \mathbf{v}_C$

[0,0,0,1,0,0,0,0,0,0]

+

[1,0,0,0,0,0,0,0,0,0]

+

[0,0,0,0,0,0,1,0,0,0]

+

[0,0,1,0,0,0,0,0,0,0]

=

[1,0,0,1,0,0,1,0,0,0]

Encoding k elements

$\text{encode}(D, A, G, C)$

A= [1,0,0,0,0,0,0,0,0,0]

B= [0,1,0,0,0,0,0,0,0,0]

C= [0,0,1,0,0,0,0,0,0,0]

D= [0,0,0,1,0,0,0,0,0,0]

E= [0,0,0,0,1,0,0,0,0,0]

F= [0,0,0,0,0,1,0,0,0,0]

G= [0,0,0,0,0,0,1,0,0,0]

H= [0,0,0,0,0,0,0,1,0,0]

I= [0,0,0,0,0,0,0,0,1,0]

J= [0,0,0,0,0,0,0,0,0,1]

$$\mathbf{v}_D + \mathbf{v}_A + \mathbf{v}_G + \mathbf{v}_C$$

[0,0,0,1,0,0,0,0,0,0]

+

[1,0,0,0,0,0,0,0,0,0]

+

[0,0,0,0,0,0,1,0,0,0]

+

[0,0,1,0,0,0,0,0,0,0]

=

[1,0,0,1,0,0,1,0,0,0]

what does this miss?

Encoding k elements

$\text{encode}(D, A, G, C)$

A= [1,0,0,0,0,0,0,0,0,0]

B= [0,1,0,0,0,0,0,0,0,0]

C= [0,0,1,0,0,0,0,0,0,0]

D= [0,0,0,1,0,0,0,0,0,0]

E= [0,0,0,0,1,0,0,0,0,0]

F= [0,0,0,0,0,1,0,0,0,0]

G= [0,0,0,0,0,0,1,0,0,0]

H= [0,0,0,0,0,0,0,1,0,0]

I= [0,0,0,0,0,0,0,0,1,0]

J= [0,0,0,0,0,0,0,0,0,1]

$\mathbf{v}_D \circ \mathbf{v}_A \circ \mathbf{v}_G \circ \mathbf{v}_C$

Encoding k elements

$\text{encode}(D, A, G, C)$

$$\mathbf{v}_D \circ \mathbf{v}_A \circ \mathbf{v}_G \circ \mathbf{v}_C$$

$[0,0,0,1,0,0,0,0,0,0]$

\circ

$[1,0,0,0,0,0,0,0,0,0]$

\circ

$[0,0,0,0,0,0,1,0,0,0]$

\circ

$[0,0,1,0,0,0,0,0,0,0]$


$=$

$[0,0,0,1,0,0,0,0,0,1,0,0,0,0,0,0,0,0,0,0,0,0,1,0,0,0,0,0,1,0,0,0,0,0,0]$

|

|

|



$$(\mathbf{v}_D + \mathbf{v}_A + \mathbf{v}_G + \mathbf{v}_C)\mathbf{W}$$

W

$$\begin{aligned} &[0,0,0,1,0,0,0,0,0,0] \\ &\quad + \\ &[1,0,0,0,0,0,0,0,0,0] \\ &\quad + \\ &[0,0,0,0,0,0,1,0,0,0] \\ &\quad + \\ &[0,0,1,0,0,0,0,0,0,0] \\ &\quad = \\ &\mathbf{[1,0,0,1,0,0,1,0,0,0]} \end{aligned}$$

A=	[-0.32, 0.09, 0.33,-0.44]
B=	[0.29, 0.02,-0.46,-0.39]
C=	[-0.46, 0.24,-0.16, 0.08]
D=	[-0.15,-0.31, 0.34, 0.00]
E=	[-0.10,-0.37, 0.01, 0.40]
F=	[-0.28,-0.26,-0.24, 0.31]
G=	[-0.32,-0.42,-0.21, 0.18]
H=	[-0.09,-0.01, 0.06, 0.14]
I=	[0.28,-0.02,-0.39, 0.12]
J=	[0.23,-0.22,-0.14, 0.28]

$$(\mathbf{v}_D + \mathbf{v}_A + \mathbf{v}_G + \mathbf{v}_C) \mathbf{W}$$

$$= \mathbf{v}_D \cdot \mathbf{W} + \mathbf{v}_A \cdot \mathbf{W} + \mathbf{v}_G \cdot \mathbf{W} + \mathbf{v}_C \cdot \mathbf{W}$$

W

$$\begin{aligned}
 &[0,0,0,1,0,0,0,0,0,0] \\
 &+ \\
 &[1,0,0,0,0,0,0,0,0,0] \\
 &+ \\
 &[0,0,0,0,0,0,1,0,0,0] \\
 &+ \\
 &[0,0,1,0,0,0,0,0,0,0] \\
 &= \\
 &\mathbf{[1,0,0,1,0,0,1,0,0,0]}
 \end{aligned}$$

$$\begin{aligned}
 A &= [-0.32, 0.09, 0.33, -0.44] \\
 B &= [0.29, 0.02, -0.46, -0.39] \\
 C &= [-0.46, 0.24, -0.16, 0.08] \\
 D &= [-0.15, -0.31, 0.34, 0.00] \\
 E &= [-0.10, -0.37, 0.01, 0.40] \\
 F &= [-0.28, -0.26, -0.24, 0.31] \\
 G &= [-0.32, -0.42, -0.21, 0.18] \\
 H &= [-0.09, -0.01, 0.06, 0.14] \\
 I &= [0.28, -0.02, -0.39, 0.12] \\
 J &= [0.23, -0.22, -0.14, 0.28]
 \end{aligned}$$

$$\begin{aligned}
 & (\mathbf{v}_D + \mathbf{v}_A + \mathbf{v}_G + \mathbf{v}_C) \mathbf{W} \\
 &= \mathbf{v}_D \cdot \mathbf{W} + \mathbf{v}_A \cdot \mathbf{W} + \mathbf{v}_G \cdot \mathbf{W} + \mathbf{v}_C \cdot \mathbf{W}
 \end{aligned}$$

sum of rows in **W**

each row corresponds to a certain vocabulary item.

$$(\mathbf{v}_D \circ \mathbf{v}_A \circ \mathbf{v}_G \circ \mathbf{v}_C) \mathbf{W}$$

$$(\mathbf{v}_D \circ \mathbf{v}_A \circ \mathbf{v}_G \circ \mathbf{v}_C) \mathbf{W}$$

still sum of rows in **W**
but **W** has 4x many rows.

$$(\mathbf{v}_D \circ \mathbf{v}_A \circ \mathbf{v}_G \circ \mathbf{v}_C) \mathbf{W}$$

still sum of rows in **W**
but **W** has 4x many rows.

alternatively:

$$= \mathbf{v}_D \cdot \mathbf{W}' + \mathbf{v}_A \cdot \mathbf{W}'' + \mathbf{v}_G \cdot \mathbf{W}''' + \mathbf{v}_C \cdot \mathbf{W}''''$$

$$(\mathbf{v}_D \circ \mathbf{v}_A \circ \mathbf{v}_G \circ \mathbf{v}_C) \mathbf{W}$$

still sum of rows in **W**
but **W** has 4x many rows.

alternatively:

$$= \mathbf{v}_D \cdot \mathbf{W}' + \mathbf{v}_A \cdot \mathbf{W}'' + \mathbf{v}_G \cdot \mathbf{W}''' + \mathbf{v}_C \cdot \mathbf{W}''''$$

$$\mathbf{W} = \mathbf{W}' \circ \mathbf{W}'' \circ \mathbf{W}''' \circ \mathbf{W}''''$$

[0,0,0,1,0,0,0,0,0,0]

o

[1,0,0,0,0,0,0,0,0,0]

o

[0,0,0,0,0,0,1,0,0,0]

o

[0,0,1,0,0,0,0,0,0,0]

A(-3)=	[0.42,-0.15, 0.12, 0.02]
B(-3)=	[0.28,-0.15,-0.11, 0.32]
C(-3)=	[0.15,-0.24, 0.23, 0.41]
D(-3)=	[-0.12,-0.24, 0.12,-0.34]
E(-3)=	[-0.42,-0.21, 0.08, 0.40]
F(-3)=	[0.20, 0.11,-0.31, 0.33]
G(-3)=	[0.07,-0.05, 0.16, 0.23]
H(-3)=	[0.28, 0.03, 0.22,-0.49]
I(-3)=	[0.08, 0.39,-0.25, 0.27]
J(-3)=	[0.10,-0.42,-0.37, 0.35]
<hr/>	
A(-2)=	[-0.00, 0.41, 0.19, 0.49]
B(-2)=	[0.24, 0.48, 0.34,-0.42]
C(-2)=	[-0.46, 0.22, 0.24,-0.21]
D(-2)=	[-0.11,-0.48, 0.18,-0.22]
E(-2)=	[-0.32, 0.10,-0.41,-0.43]
F(-2)=	[0.32, 0.02,-0.22, 0.06]
G(-2)=	[-0.31,-0.36, 0.09, 0.39]
H(-2)=	[0.01,-0.22,-0.09,-0.15]
I(-2)=	[0.01, 0.10,-0.16,-0.21]
J(-2)=	[-0.24, 0.40,-0.34,-0.13]
<hr/>	
A(-1)=	[-0.23,-0.38, 0.02, 0.32]
B(-1)=	[-0.34, 0.04,-0.18,-0.00]
C(-1)=	[0.40,-0.02, 0.10,-0.16]
D(-1)=	[0.13,-0.07,-0.19,-0.01]
E(-1)=	[0.40, 0.27,-0.33, 0.36]
F(-1)=	[0.04,-0.13,-0.43, 0.39]
G(-1)=	[0.44, 0.38, 0.03,-0.39]
H(-1)=	[0.41,-0.23, 0.33,-0.08]
I(-1)=	[-0.50,-0.16,-0.42,-0.27]
J(-1)=	[-0.15, 0.41, 0.46,-0.16]
<hr/>	
A(+0)=	[-0.11, 0.03, 0.20, 0.50]
B(+0)=	[0.16,-0.34, 0.20,-0.21]
C(+0)=	[0.05,-0.13,-0.23,-0.31]
D(+0)=	[0.13,-0.02, 0.38,-0.09]
E(+0)=	[0.30, 0.39, 0.10, 0.38]
F(+0)=	[-0.16,-0.31,-0.02,-0.34]
G(+0)=	[0.06,-0.04, 0.02,-0.32]
H(+0)=	[0.25, 0.30, 0.29, 0.24]
I(+0)=	[0.40,-0.17,-0.18,-0.19]
J(+0)=	[0.27, 0.33,-0.42,-0.07]

W'

W''

W'''

W''''

- 1-hot times matrix: row selection
- sum of 1-hot times matrix: row selection + sum
- concat of 1-hot: like using 1-hot from larger vocab

"Embedding Layer"

- Very common in neural network land:
 - associate each vocabulary item with a row in matrix **E** of dense vectors (dim of row $\ll |V|$)
 - concat or sum rows of **E** for input.

"Embedding Layer"

$\text{encode}(D, A, G, C)$

$$= \mathbf{E}_{[D]} \circ \mathbf{E}_{[A]} \circ \mathbf{E}_{[G]} \circ \mathbf{E}_{[C]}$$

E

A= [-0.32, 0.09, 0.33,-0.44]
B= [0.29, 0.02,-0.46,-0.39]
C= [-0.46, 0.24,-0.16, 0.08]
D= [-0.15,-0.31, 0.34, 0.00]
E= [-0.10,-0.37, 0.01, 0.40]
F= [-0.28,-0.26,-0.24, 0.31]
G= [-0.32,-0.42,-0.21, 0.18]
H= [-0.09,-0.01, 0.06, 0.14]
I= [0.28,-0.02,-0.39, 0.12]
J= [0.23,-0.22,-0.14, 0.28]

[-0.15,-0.31, 0.34, 0.00,-0.32, 0.09, 0.33,-0.44,-0.32,-0.42,-0.21, 0.18,-0.46, 0.24,-0.16, 0.08]

$$\begin{array}{rcc}
& & \mathbb{R}^{d_{out}} \\
\text{softmax}(\square) & & \\
\uparrow & & \mathbb{R}^{d_{out}} \\
\square \mathbf{W}^3 + \mathbf{b}^3 & & \\
\uparrow & & \mathbb{R}^{d_2} \\
g(\square \mathbf{W}^2 + \mathbf{b}^2) & & \\
\uparrow & & \mathbb{R}^{d_1} \\
g(\square \mathbf{W}^1 + \mathbf{b}^1) & & \\
\uparrow & & \mathbb{R}^{d_{in}} \\
\mathbf{E}_{[x_{k-4}]} \circ \mathbf{E}_{[x_{k-3}]} \circ \mathbf{E}_{[x_{k-2}]} \circ \mathbf{E}_{[x_{k-1}]} & & \\
\uparrow & & \\
\text{encode}(x_{k-4}, x_{k-3}, x_{k-2}, x_{k-1}) & &
\end{array}$$

"Embedding Layer"

$\text{encode}(D, A, G, C)$

E

$$= \mathbf{E}_{[D]} \circ \mathbf{E}_{[A]} \circ \mathbf{E}_{[G]} \circ \mathbf{E}_{[C]}$$

$[-0.15, -0.31, 0.34, 0.00, -0.32, 0.09, 0.33, -0.44, -0.32, -0.42, -0.21, 0.18, -0.46, 0.24, -0.16, 0.08]$

A= $[-0.32, 0.09, 0.33, -0.44]$
B= $[0.29, 0.02, -0.46, -0.39]$
C= $[-0.46, 0.24, -0.16, 0.08]$
D= $[-0.15, -0.31, 0.34, 0.00]$
E= $[-0.10, -0.37, 0.01, 0.40]$
F= $[-0.28, -0.26, -0.24, 0.31]$
G= $[-0.32, -0.42, -0.21, 0.18]$
H= $[-0.09, -0.01, 0.06, 0.14]$
I= $[0.28, -0.02, -0.39, 0.12]$
J= $[0.23, -0.22, -0.14, 0.28]$

how does this relate to what
we had before?

what is

$$(\mathbf{E}_{[D]} \circ \mathbf{E}_{[A]} \circ \mathbf{E}_{[G]} \circ \mathbf{E}_{[C]}) \mathbf{W}$$

?

"Embedding Layer"

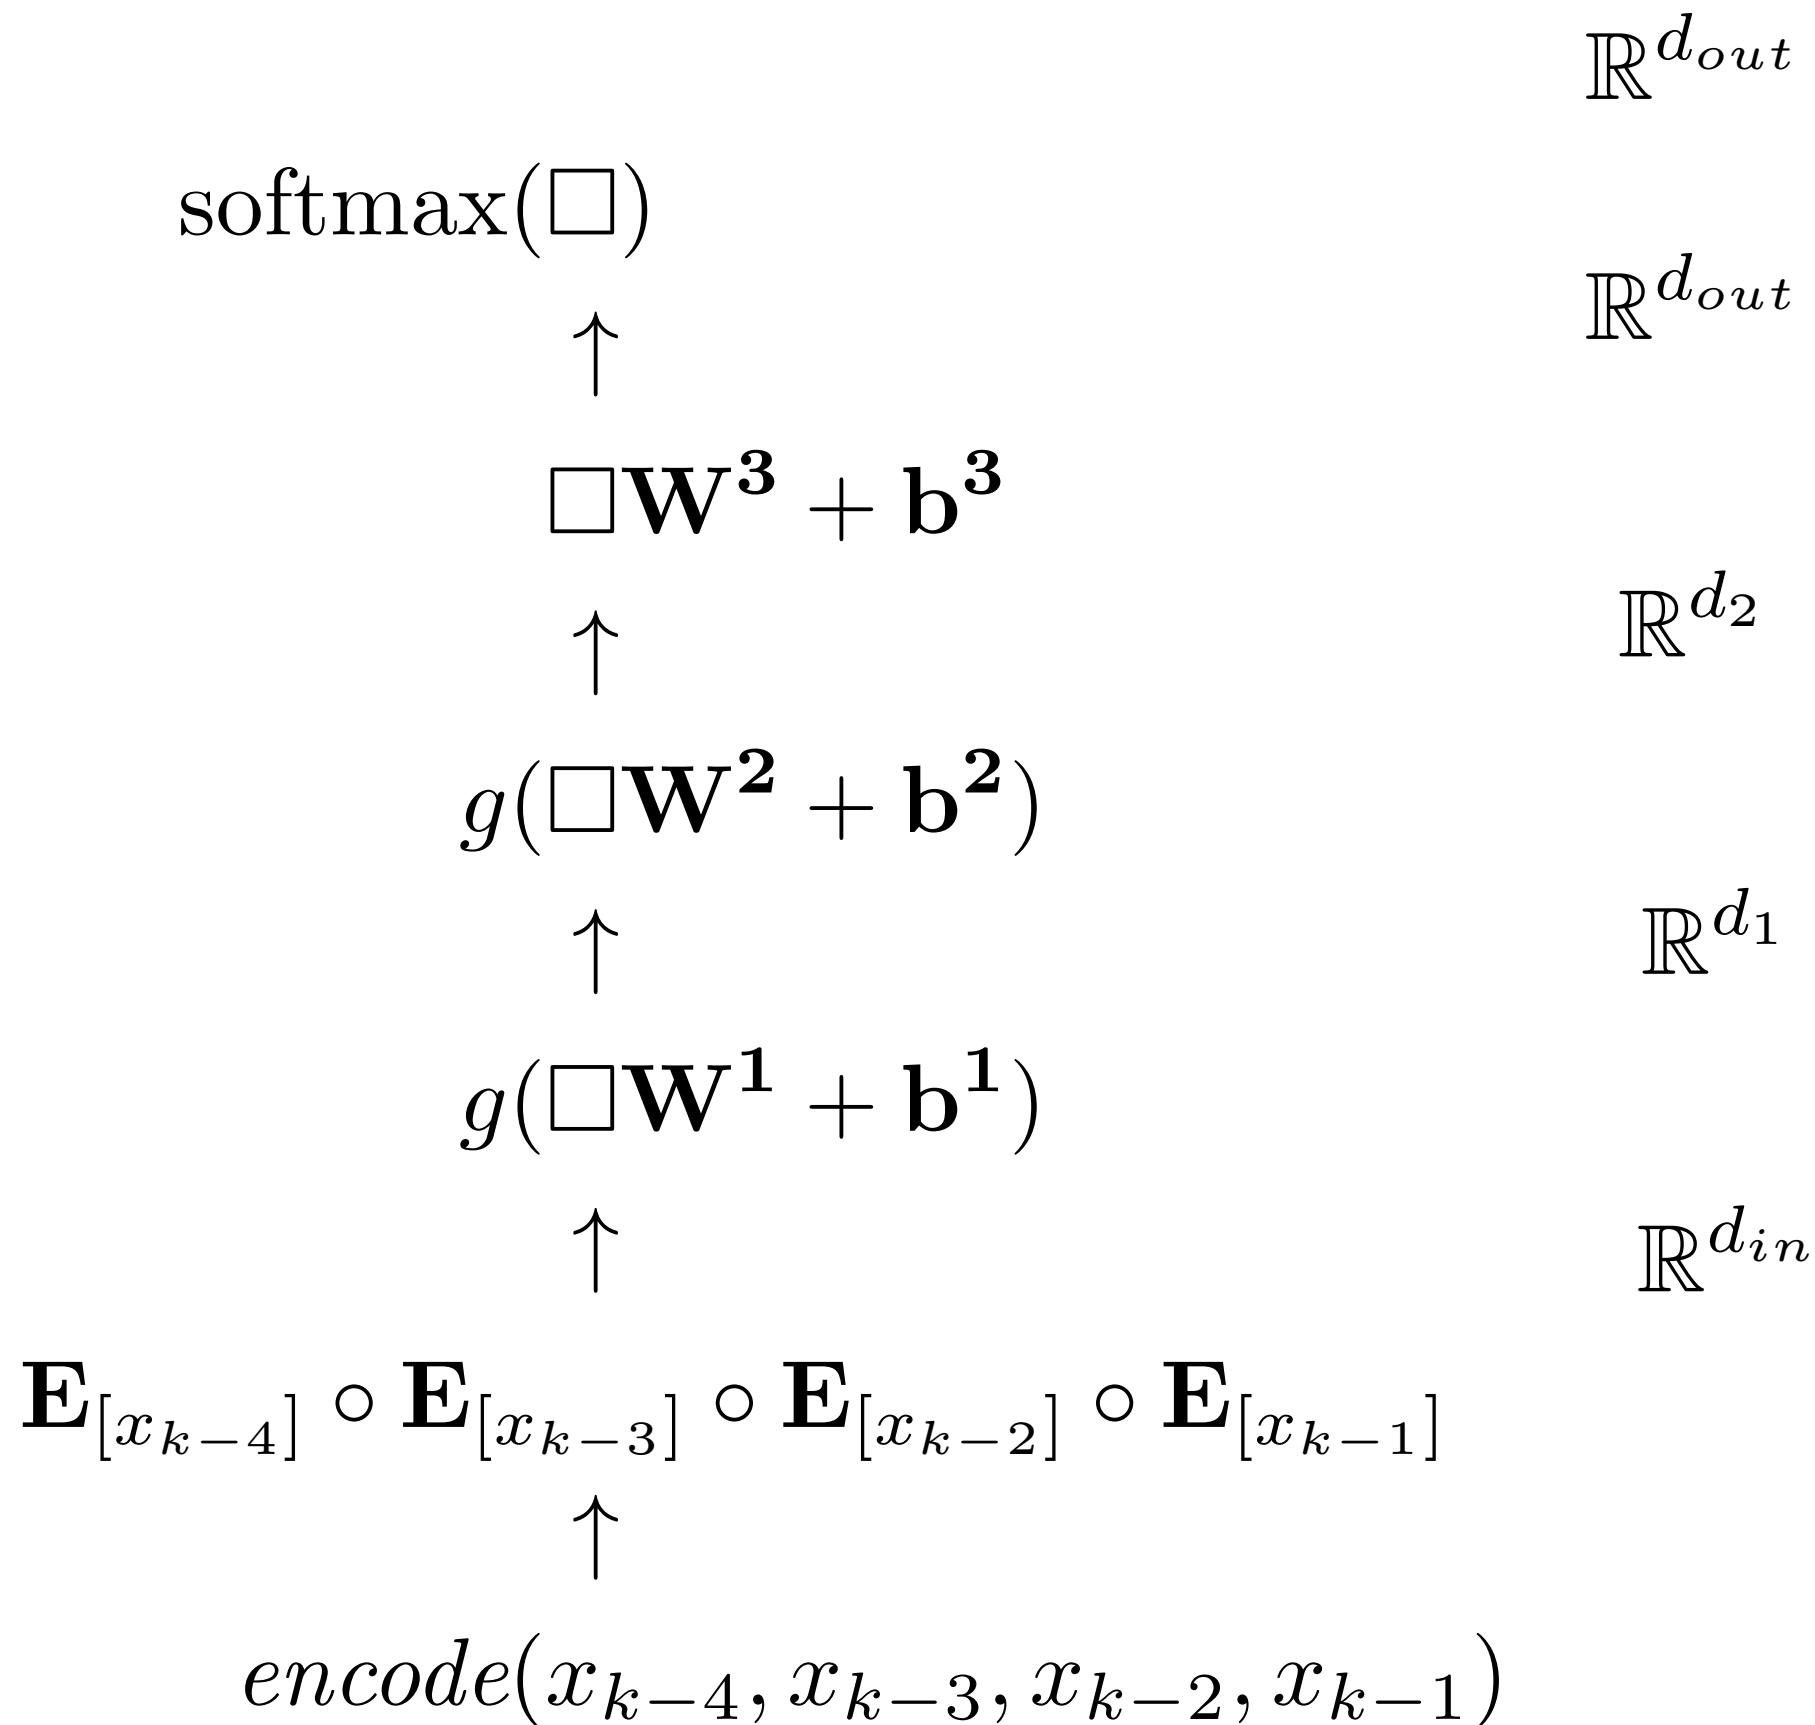
$$(\mathbf{E}_{[D]} \circ \mathbf{E}_{[A]} \circ \mathbf{E}_{[G]} \circ \mathbf{E}_{[C]}) \mathbf{W}$$

same row in \mathbf{E} regardless of position in the sequence.

but \mathbf{W} will transform this row differently for each position.

$$\begin{array}{rcc}
& & \mathbb{R}^{d_{out}} \\
\text{softmax}(\square) & & \\
\uparrow & & \mathbb{R}^{d_{out}} \\
\square \mathbf{W}^3 + \mathbf{b}^3 & & \\
\uparrow & & \mathbb{R}^{d_2} \\
g(\square \mathbf{W}^2 + \mathbf{b}^2) & & \\
\uparrow & & \mathbb{R}^{d_1} \\
g(\square \mathbf{W}^1 + \mathbf{b}^1) & & \\
\uparrow & & \mathbb{R}^{d_{in}} \\
\mathbf{E}_{[x_{k-4}]} \circ \mathbf{E}_{[x_{k-3}]} \circ \mathbf{E}_{[x_{k-2}]} \circ \mathbf{E}_{[x_{k-1}]} & & \\
\uparrow & & \\
\text{encode}(x_{k-4}, x_{k-3}, x_{k-2}, x_{k-1}) & &
\end{array}$$

what's in **W1**?
 what's in **W2**?
 what's in **W3**?



Neural LM

$\text{softmax}(\square)$

\uparrow

$$\square \mathbf{W}^3 + \mathbf{b}^3$$

\uparrow

$$g(\square \mathbf{W}^2 + \mathbf{b}^2)$$

\uparrow

$$g(\square \mathbf{W}^1 + \mathbf{b}^1)$$

\uparrow

$$\mathbf{E}_{[x_{k-4}]} \circ \mathbf{E}_{[x_{k-3}]} \circ \mathbf{E}_{[x_{k-2}]} \circ \mathbf{E}_{[x_{k-1}]}$$

\uparrow

$$\text{encode}(x_{k-4}, x_{k-3}, x_{k-2}, x_{k-1})$$

Neural LM

A Neural Probabilistic Language Model

Yoshua Bengio

Réjean Ducharme

Pascal Vincent

Christian Jauvin

Département d'Informatique et Recherche Opérationnelle

Centre de Recherche Mathématiques

Université de Montréal, Montréal, Québec, Canada

BENGIOY@IRO.UMONTREAL.CA

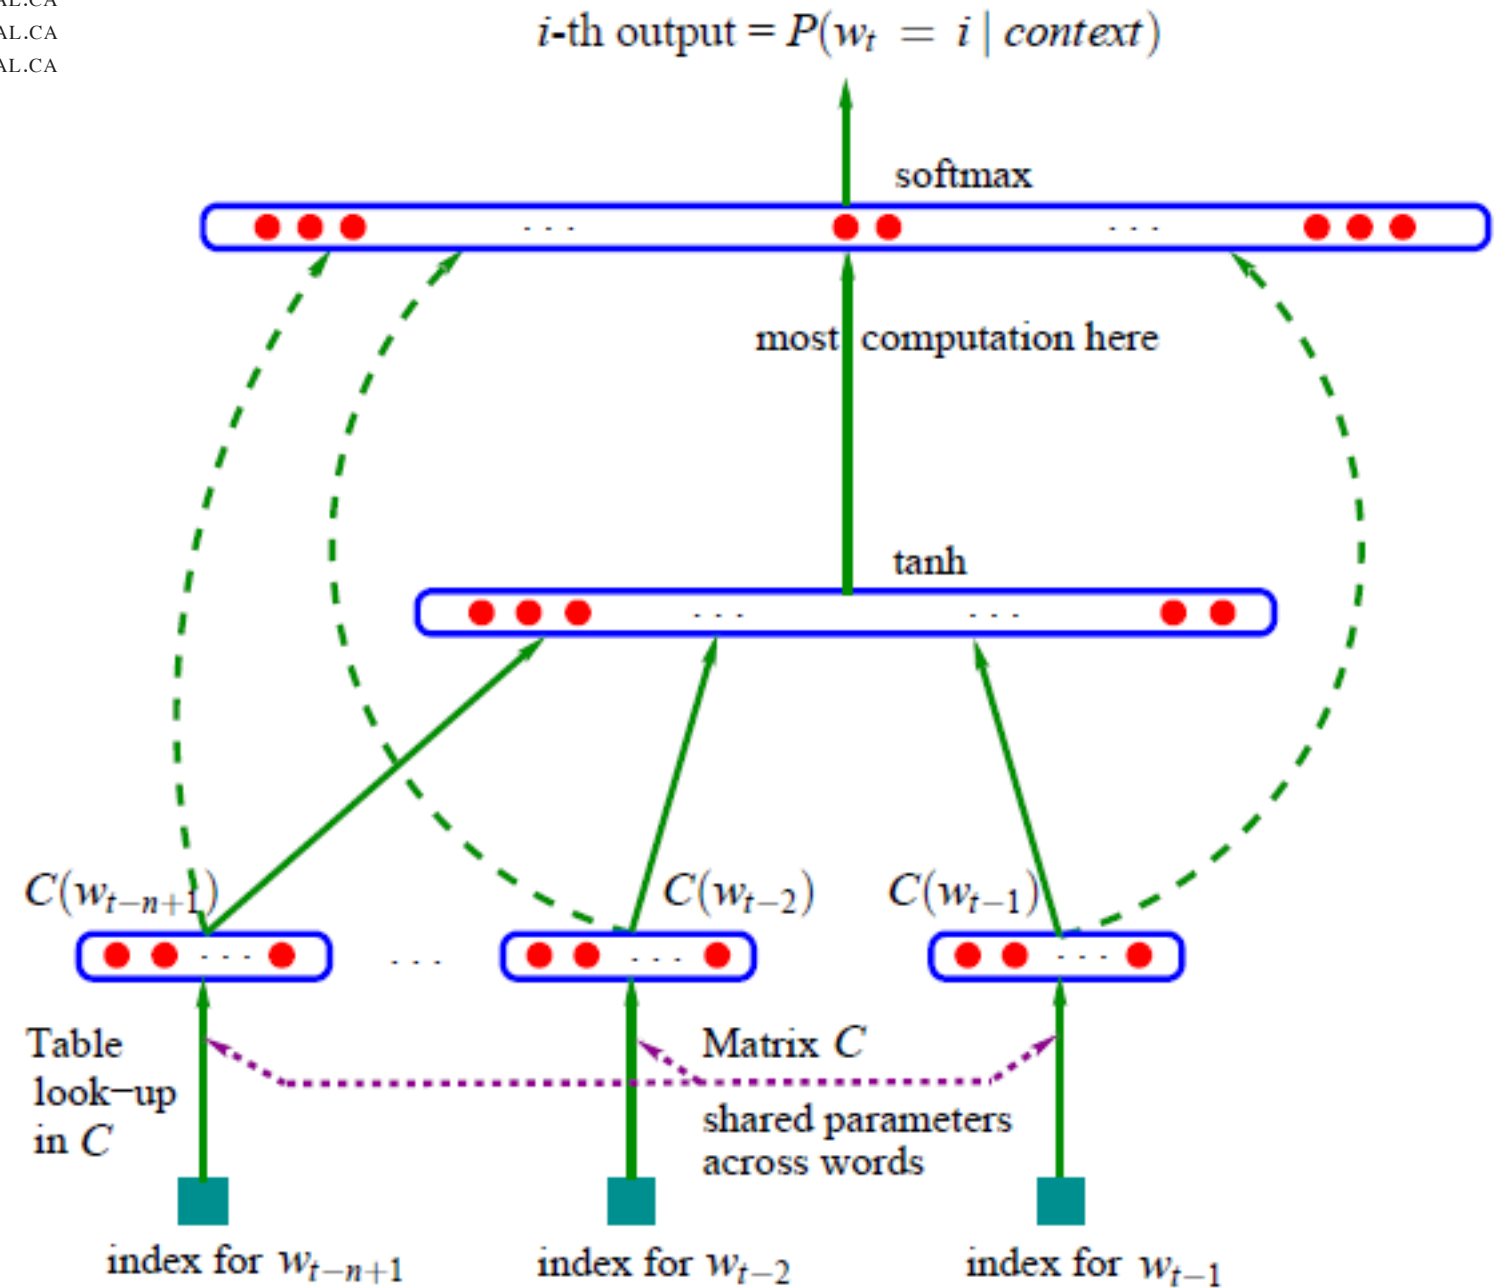
DUCHARME@IRO.UMONTREAL.CA

VINCENTP@IRO.UMONTREAL.CA

JAUVINC@IRO.UMONTREAL.CA

Bengio et al 2003

Pretty much same model,
But fewer layers.



Neural LM

- What is the cost of increasing the history length?

Neural LM

- We can use a trained LM for scoring a given sentence. (how?)

Neural LM

- We can use a trained LM for comparing two given sentences. (how?)

Neural LM

- We can use a trained LM for **generating new sentences**. (how?)

Neural LM

- We can use K trained LMs for **k-class classification**. (how?)

Neural LM

- We can use K trained LMs for **k-class classification**. (how?)

$$p(w_1, w_1, \dots, w_n | LM_1)$$

$$p(w_1, w_1, \dots, w_n | LM_2)$$

$$p(w_1, w_1, \dots, w_n | LM_3)$$

$$\hat{y} = \arg \max_k p(w_1, w_1, \dots, w_n | LM_k)$$

sequence prediction tasks

- In LM:
 - predict **word i** based on **k previous words**.
- But we could also predict **label** based on **k items**.
 - which tasks can this be used for?

classification



ACGCGCTCGATCG

V

GTGCGCTCGAACG

V

ACGCGTTCTACCG

X

tagging

DET	ADJ	NOUN	VERB	PREP	DET	ADJ	NOUN
The	brown	fox	jumped	over	the	lazy	dog

tagging

Window-based approach

NOUN



The brown **fox** jumped over

VERB

brown fox **jumped** over the

PREP

fox jumped **over** the lazy

back to LM

Training a language model

- Set dimensions of **E**, **W3**, according to vocab size.
- Initial random values for **E**, **W1**, **W2**, **W3**, **b1**, **b2**, **b3**
- For every n-tuple in some text:
 - try to predict last item based on prev n-1
 - use cross-entropy loss.

What happens after training?

- Consider the columns of **W_3** .
- Consider the rows of **E** .

$\mathbb{R}^{d_{out}}$ $\text{softmax}(\square)$ $\mathbb{R}^{d_{out}}$ \uparrow

$$\square \mathbf{W}^3 + \mathbf{b}^3$$

what's in **W3**?

 \uparrow \mathbb{R}^{d_2}

$$g(\square \mathbf{W}^2 + \mathbf{b}^2)$$

 \uparrow \mathbb{R}^{d_1}

$$g(\square \mathbf{W}^1 + \mathbf{b}^1)$$


 \uparrow $\mathbb{R}^{d_{in}}$

$$\mathbf{E}_{[x_{k-4}]} \circ \mathbf{E}_{[x_{k-3}]} \circ \mathbf{E}_{[x_{k-2}]} \circ \mathbf{E}_{[x_{k-1}]}$$

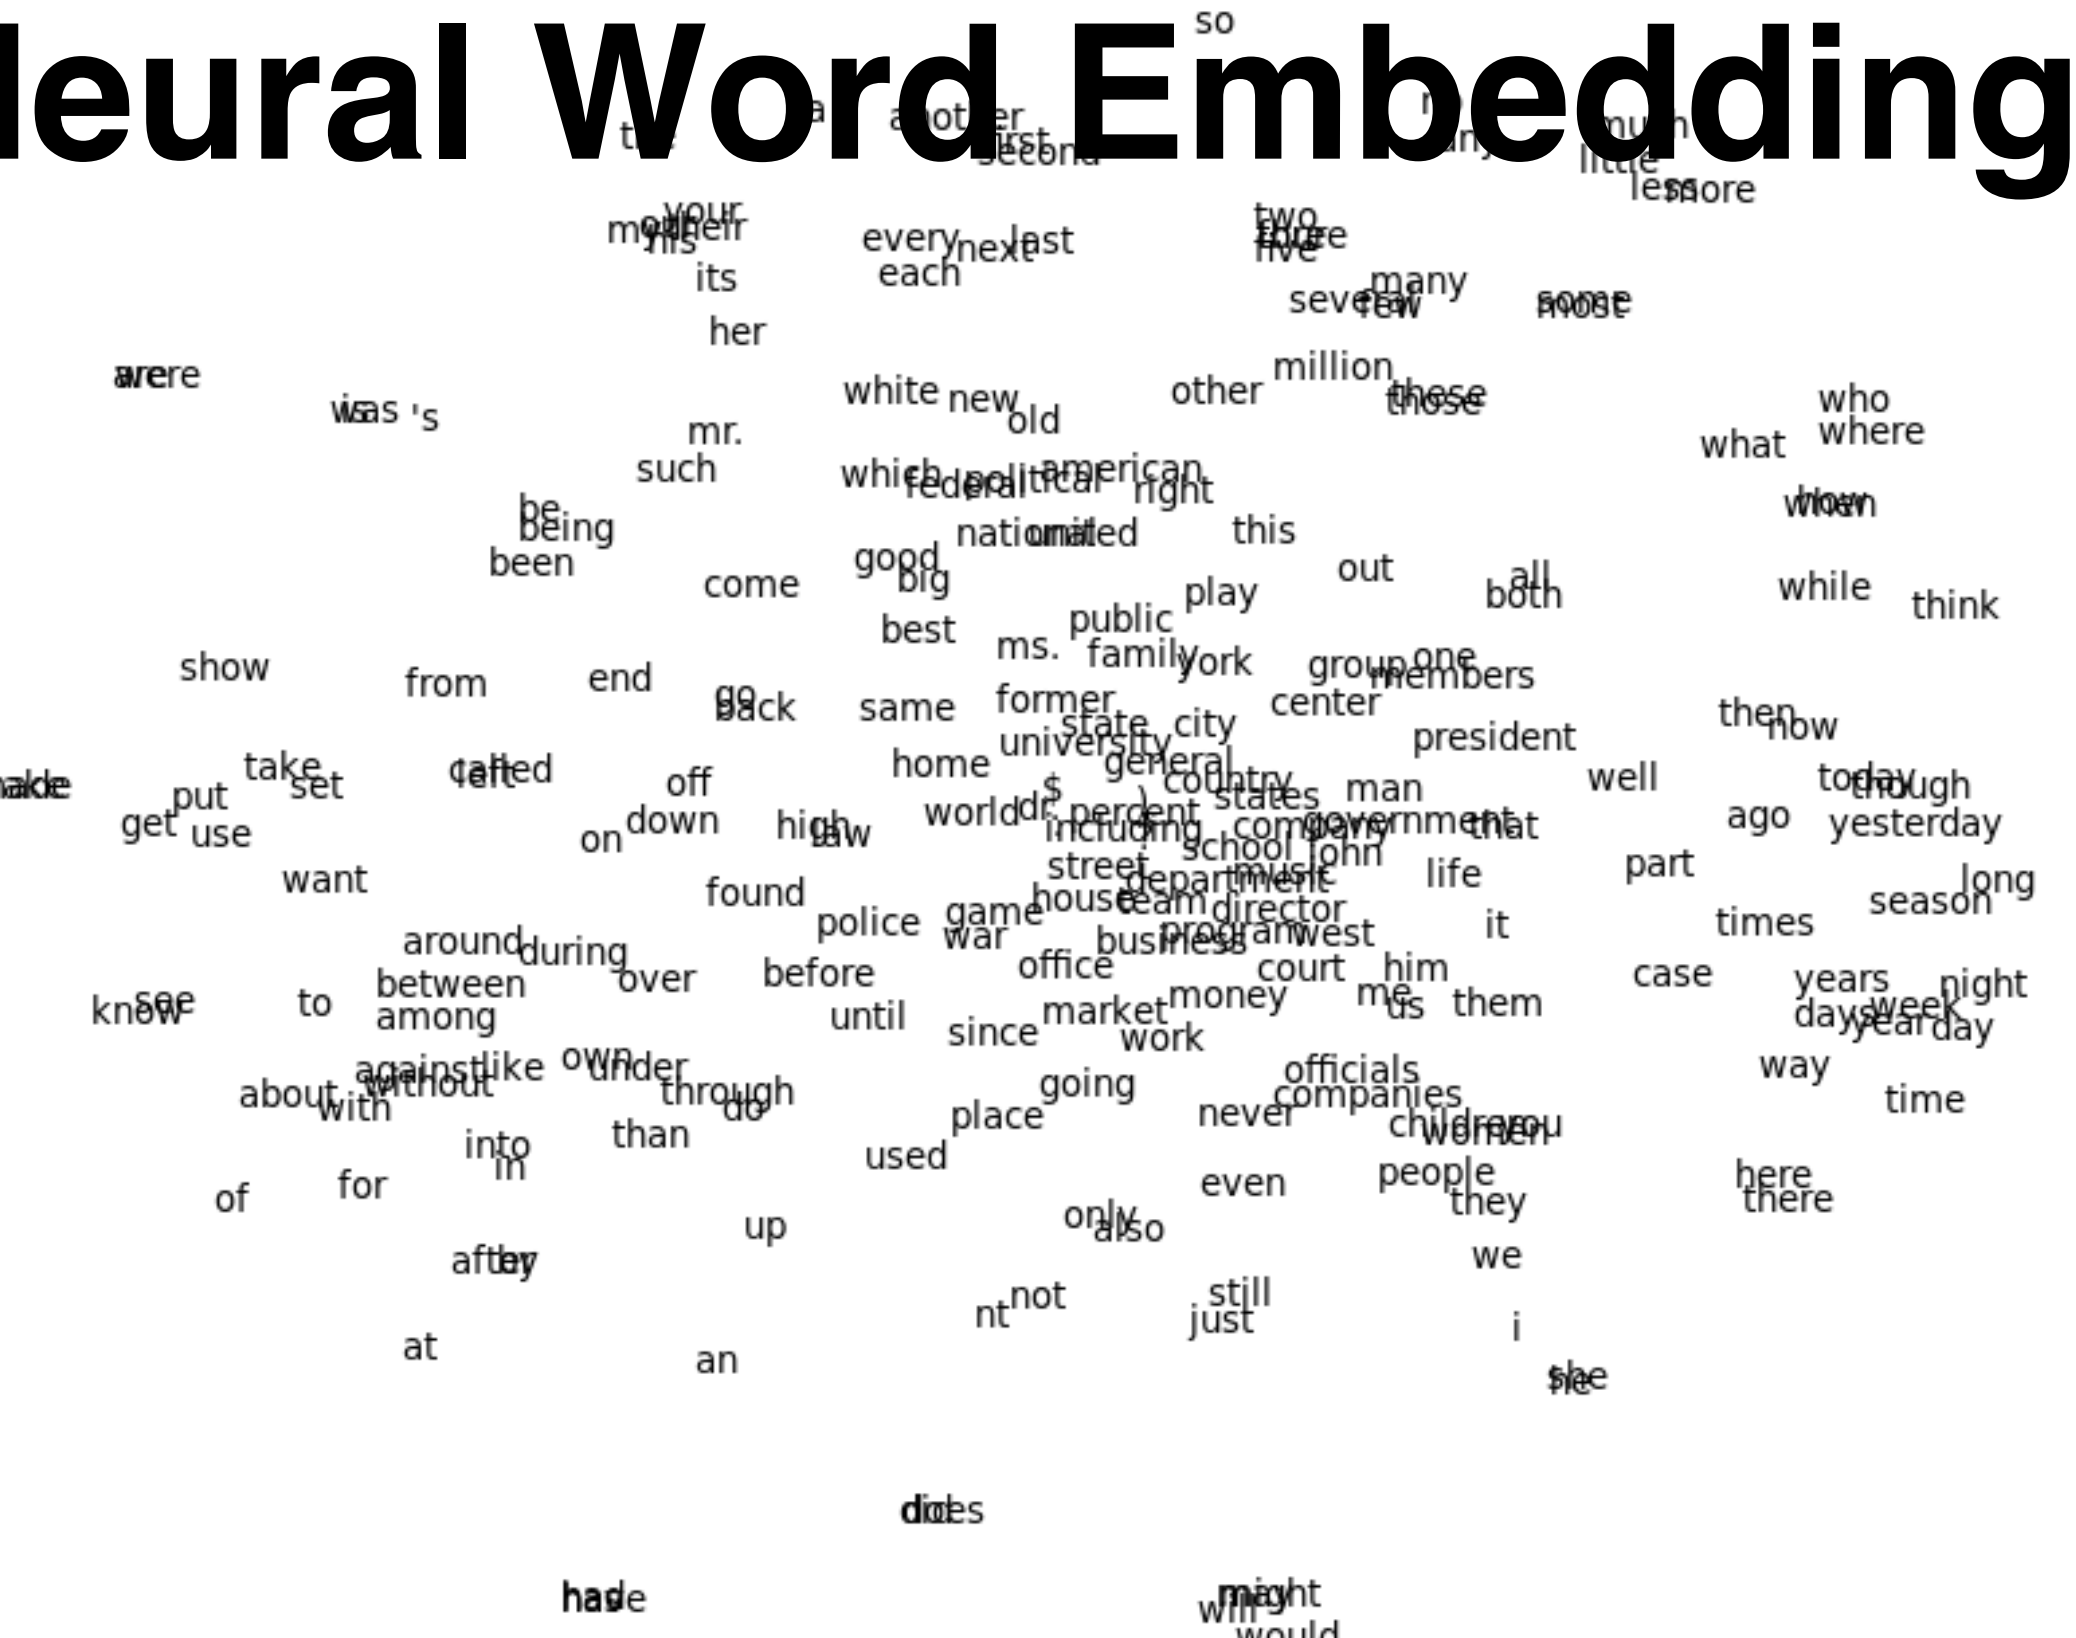
 \uparrow

$$\text{encode}(x_{k-4}, x_{k-3}, x_{k-2}, x_{k-1})$$

columns of **W3**
correspond
to vocab items!

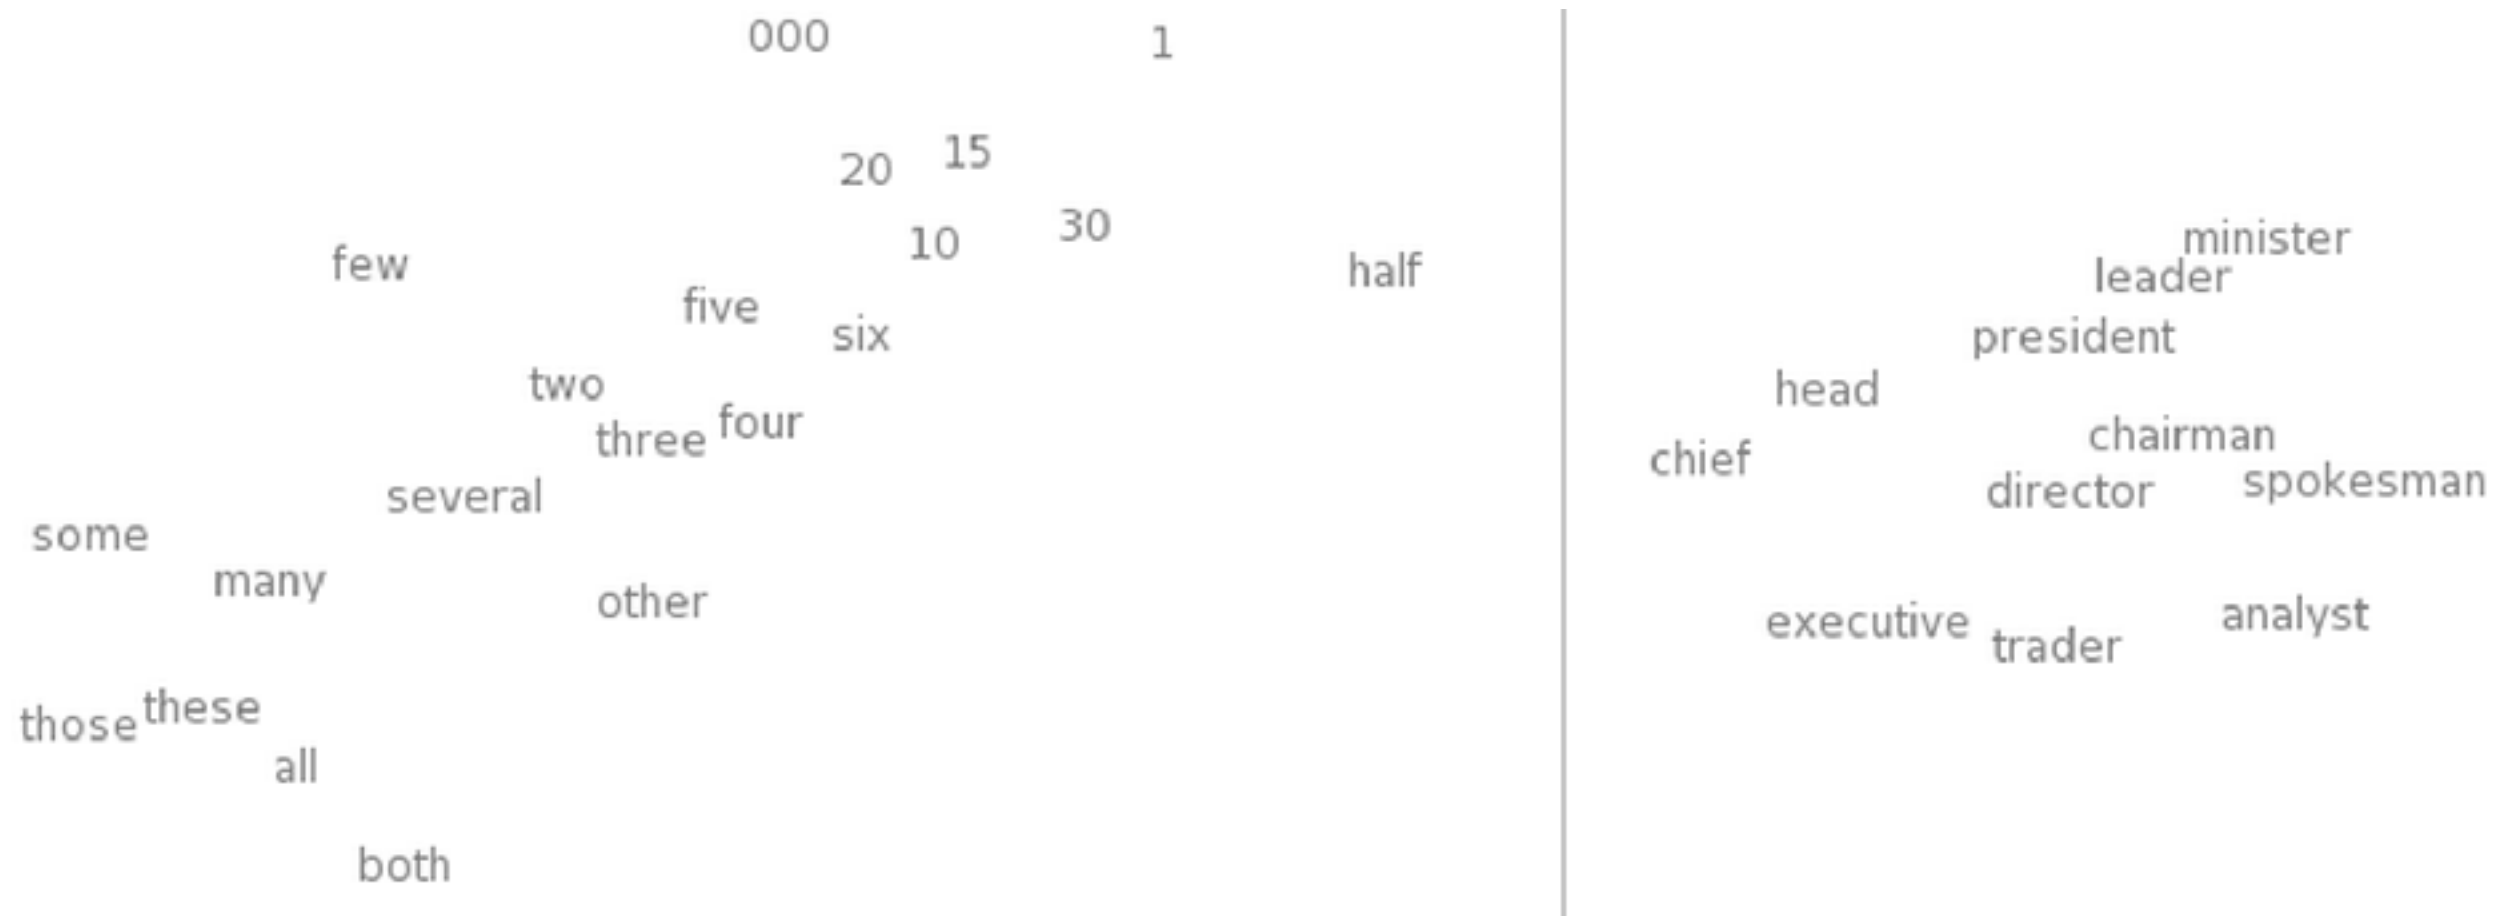


Neural Word Embeddings




The plot displays a large number of words as points in a 2D space. The x-axis and y-axis both range from -15 to 15, with major grid lines every 5 units. The words are distributed across the plot, with some appearing more frequently or in larger fonts than others. The words are not arranged in any specific order, but they represent the output of a neural network trained on a large corpus of text. The plot illustrates the concept of word embeddings, where words are represented as vectors in a high-dimensional space.

Neural Word Embeddings



Neural Word Embeddings

Country and Capital Vectors Projected by PCA



Neural Word Embeddings

Target Word	BoW5
batman	nightwing aquaman catwoman superman manhunter
hogwarts	dumbledore hallows half-blood malfoy snape
turing	nondeterministic non-deterministic computability deterministic finite-state
florida	gainesville fla jacksonville tampa lauderdale
object-oriented	aspect-oriented smalltalk event-driven prolog domain-specific
dancing	singing dance dances dancers tap-dancing

Neural Word Embeddings

- We trained a language model.
- We ended up with **vector representations of words**.
- These representations are useful -- they encode various aspects of word similarity.

tagging + pre-training

we can use the **E** we got from LM training to initialize **E** for the POS tagging task.

(why is that helpful?)

NOUN

The brown **fox** jumped over

VERB

brown fox **jumped** over the

PREP

fox jumped **over** the lazy

pre-training

- This is a sort of **semi-supervised** learning or **multi-task** learning.
- We learn from "unannotated" data.
- We then use the representations on tasks with annotated data.

Neural Word Embeddings

- We trained a language model.
- We ended up with **vector representations of words**.
- These representations are useful -- they encode various aspects of word similarity.
- **A form of semi-supervised learning.**
- More next week.